

XV6 Lottery Scheduler

Team Members

Dagan Martinez

Source Added and Changed

xv6 root

Makefile.c

⚙	@@ -22,11 +22,15 @@ all: xv6.img fs.img	
22	22	# http://gcc.gnu.org/onlinedocs/gcc-4.4.6/gcc/Invoking-GCC.html
23	23	CC = gcc
24	24	# enable extra warnings
25		-CFLAGS += -Wall -Wno-deprecated-declarations
	25	+CFLAGS += -Wall -Wextra
	26	+# Disable some warnings
	27	+CFLAGS += -Wno-deprecated-declarations -Wno-sign-compare -Wno-unused-parameter -Wno-implicit-fallthrough
26	28	# treat warnings as errors
27	29	CFLAGS += -Werror
28	30	# produce debugging information for use by gdb
29	31	CFLAGS += -ggdb
	32	+# Use a modern version of C
	33	+CFLAGS += -std=gnu99
30	34	
31	35	# uncomment to enable optimizations. improves performance, but may make
32	36	# debugging more difficult
⚙	@@ -80,7 +84,7 @@ QEMUGDB := \$(shell if \$(QEMU) -help grep -q '^gdb'; \	
80	84	
81	85	# number of CPUs to emulate in QEMU
82	86	ifndef CPUS
83		-CPUS := 2
	87	+CPUS := 1
84	88	endif
85	89	
86	90	QEMUOPTS := -hdb fs.img xv6.img -smp \$(CPUS)
⚙		

I changed the number of CPUs to be run from 2 to 1, to better show lottery scheduling with two processes. I also added and suppressed some errors to taste, and fixed the C version instead of falling back on the default choice of that version of GCC.

xv6/include

Random.h

```
1 #ifndef _RANDOM_H
2 #define _RANDOM_H
3 /*
4  * Multiply with carry code by George Marsaglia
5  *
6  * Modified for XV6 by Dagan Martinez
7  */
8 #include "types.h"
9
10 #define PHI 0x9e3779b9
11
12 static uint Q[4096], c = 362436;
13
14 static void srand(uint x)
15 {
16     int i;
17
18     Q[0] = x;
19     Q[1] = x + PHI;
20     Q[2] = x + PHI + PHI;
21
22     for (i = 3; i < 4096; i++)
23         Q[i] = Q[i - 3] ^ Q[i - 2] ^ PHI ^
24         i;
25
26 static uint rand(void)
27 {
28     if(sizeof(unsigned long long) != 8){
29         return 0;
30     }
31     unsigned long long t, a = 18782LL;
32     static uint i = 4095;
33     uint x, r = 0xffffffff;
34     i = (i + 1) & 4095;
35     t = a * Q[i] + c;
36     c = (t >> 32);
37     x = t + c;
38     if (x < c) {
39         x++;
40         c++;
41     }
42     return (Q[i] = r - x);
43 }
44 #endif
```

include/random.h 1,1 Top include/random.h 44,1 Bot

Most of this code was taken from Wikipedia, with some slight modification. I put static functions in a header file so the same code would be available to both the kernel and user(for easier testing).

Pstat.h

```
14 #ifndef _PSTAT_H the process is running
13 #define _PSTAT_H
12
11 #include "param.h"
10
9 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
8
7 // Copied directly from the assignment sheet
6 // Everything in this structure is DESCRIPTIVE
5 // but NOT PRESCRIPTIVE
4 //
3 // Meaning that the values of the fields here
2 // should have absolutely no effect on the kernel
1 // (but maybe the userspace if the programmer is being an idiot)
0 struct pstat {
1     // Whether the process is running
2     // Each entry will be 1 or 0
3     //
4     // There should only be at most one '1'
5     // (Per CPU)
6     _Bool inuse[NPROC];
7
8     // PID of each process
9     int pid[NPROC];
10
11     // Number of ticks each process has accumulated
12     int ticks[NPROC];
13
14     // Number of tickets
15     int tickets[NPROC];
16
17     //state
18     enum procstate state[NPROC];
19
20     int total_tickets;
21
22 };
23
24 #endif
```

include/pstat.h

15,5

All

This structure is purely descriptive. In addition to the three fields given in Miniproject2.pdf, I added the tickets and state fields, as well as a non-array field: total_tickets, which holds the amount of ticks held by all processes combined(not counting sleeping processes)

Syscall.h

```
26  #define SYS_getpinfo 22
27  #define SYS_settickets 23
```

I had to add the syscall numbers for my two new syscalls

xv6/kernel

proc.c

```
7      -#include "spinlock.h"
      7  #include "random.h"
      8  #define STORE_TICKETS_ON_SLEEP
8      9
9      -struct {
10     -   struct spinlock lock;
11     -   struct proc proc[NPROC];
12     -} ptable;
      10  +struct ptable_type ptable = {0};
```

I had to move the struct definition of ptable to proc.c, so I could access ptable in sysproc.c (for my getpinfo syscall)

STORE_TICKETS_ON_SLEEP is an option to not count sleeping processes in the lottery

```

26  +// Keep track of the amount of tickets handed out
27  +int total_tickets;
28  +// This function should always be in a lock
29  +void setproctickets(struct proc* pp, int n)
30  +{
31  +    total_tickets -= pp->tickets;
32  +    pp->tickets = n;
33  +    total_tickets += pp->tickets;
34  +}
35  +
36  +// Just after sleeping
37  +void storetickets(struct proc* pp)
38  +{
39  +    if(pp->state != SLEEPING)
40  +        panic("Not sleeping at storetickets");
41  +#ifdef STORE_TICKETS_ON_SLEEP
42  +    total_tickets -= pp->tickets;
43  +#endif
44  +}
45  +
46  +// Just before waking
47  +void restoretickets(struct proc* pp)
48  +{
49  +    if(pp->state != SLEEPING)
50  +        panic("Not sleeping at waketickeets");
51  +#ifdef STORE_TICKETS_ON_SLEEP
52  +    total_tickets += pp->tickets;
53  +#endif
54  +}

```

These are functions used to set, store, and restore the tickets of a process

```

177  +    // Tickets
178  +    // A child will have the same number of tickets as its parent
179  +    setproctickets(np, proc->tickets);

```

When a process fork(s), its child should have the same number of tickets

```

232  +    // Remove from lottery
233  +    setproctickets(proc, 0);

```

In exit(), we have to remove the process from the lottery

```

257 +         if(p->state == ZOMBIE){
258 +             // Found one.
259 +             pid = p->pid;
260 +             kfree(p->kstack);
261 +             p->kstack = 0;
262 +             freevm(p->pgdir);
263 +             p->state = UNUSED;
264 +             p->pid = 0;
265 +             p->parent = 0;
266 +             p->name[0] = 0;
267 +             p->killed = 0;
268 +             p->ticks = 0;
269 +             setproctickets(p, 0);
270 +             release(&ptable.lock);
271 +             return pid;
272 +         }

```

When we're cleaning up zombie processes, we should also remove them from the lottery (in `wait()`) and clear the ticks

<pre> 93 void scheduler(void) 94 { 95 struct proc *p; 96 97 // Set init's tickets to 1 98 acquire(&ptable.lock); 99 setproctickets(ptable.proc, 1); 00 release(&ptable.lock); 01 02 // Seed random 03 static _Bool have_seeded = 0; 04 const int seed = 1323; 05 if(!have_seeded) 06 { 07 srand(seed); 08 have_seeded = 1; 09 } 10 11 for(;;){ 12 // Enable interrupts on this processor. 13 sti(); 14 15 // Winning ticket 16 const int golden_ticket = 17 rand()%(total_tickets + 1); 18 int ticket_count = 0; 19 20 // Loop over process table looking for process to 21 run. acquire(&ptable.lock); </pre>	<pre> 322 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) 323 { 324 // We're only looking for runnable processes 325 if(p->state != RUNNABLE) 326 { 327 #ifndef STORE_TICKETS_ON_SLEEP 328 ticket_count += p->tickets; 329 #endif 330 continue; 331 } 332 333 // Ones that have the golden ticket 334 ticket_count += p->tickets; 335 if(ticket_count < golden_ticket) 336 { 337 continue; 338 } 339 else if(ticket_count > total_tickets) 340 cprintf("Extra: %d %d %d\n", ticket_ count, total_tickets, golden_ticket); 341 342 // Switch to chosen process. It is the proc 343 // ess's job 344 // to release ptable.lock and then reacquire 345 it 346 // before jumping back to us. 347 proc = p; 348 switchvm(p); </pre>
roc.c 294,1 57%	proc.c 329,6 62%

```

345         // before jumping back to us.
346         proc = p;
347         switchvm(p);
348         p->state = RUNNING;
349         // Start timing
350         p->inuse = 1;
351         const int tickstart = ticks;
352         // Actually run process
353         swtch(&cpu->scheduler, proc->context);
354         // Record ticks
355         p->ticks += ticks - tickstart;
356         p->inuse = 0;
357
358         switchkvm();
359
360         // Process is done running for now.
361         // It should have changed its p->state before coming back.
362         proc = 0;
363         break;
364     }
365     release(&ptable.lock);
366
367 }
368 }
369

```

proc.c 369,0-1 66%

The scheduler function had the most changes. First, I had to set the first process's tickets to 1(which will set all children's processes to 1 via fork()). Next I had to seed the random function with an appropriate value.

In each iteration of the forever loop, I had to create a new winning ticket, and set the “ticket count” to zero.

When I’m foreaching through the process table, I have to decide whether or not I want to count sleeping processes in the lottery. As I defined STORE_TICKETS_ON_SLEEP previously, I do NOT add tickets from non-runnable processes.

Now, I have to filter any process that does not have the winning ticket to go back to the beginning of the loop. Note that only one process will win, after which the loop is restarted.

Before the processes is run, the timer must be resumed. After the process has run, the timer is stopped again.

```
437 // Go to sleep.
438 proc->chan = chan;
439 proc->state = SLEEPING;
440 storetickets(proc);
441 sched();
```

In sleep(), I have to store the tickets as previously mentioned. This temporarily removes them from the lottery.

```
453 // Wake up all processes sleeping on chan.
454 // The ptable lock must be held.
455 static void
456 wakeup1(void *chan)
457 {
458     struct proc *p;
459
460     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
461         if(p->state == SLEEPING && p->chan == chan)
462         {
463             restoretickets(p);
464             p->state = RUNNABLE;
465         }
466 }
```

Inside wakeup1, I restore those tickets

```

477 // Kill the process with the given pid.
478 // Process won't exit until it returns
479 // to user space (see trap in trap.c).
480     int
481 kill(int pid)
482 {
483     struct proc *p;
484
485     acquire(&ptable.lock);
486     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
487         if(p->pid == pid){
488             p->killed = 1;
489             // Wake process from sleep if necessary.
490             if(p->state == SLEEPING)
491             {
492                 restoretickets(p);
493                 p->state = RUNNABLE;
494             }
495             release(&ptable.lock);
496             return 0;
497         }
498     }
499     release(&ptable.lock);
500     return -1;
501 }
proc.c 479,2 92%

```

Kill does not actually remove the process from the ptable, so I do restore tickets if the process is sleeping(which will then almost immediately be removed again)

Proc.h

```

14 #include "pstat.h"
15 #include "spinlock.h"

```

```

65
66 // Per-process state many ticks has accumulated
67 struct proc {
68     uint sz;                // Size of process memory (bytes)
69     pde_t* pgdir;          // Page table
70     char *kstack;          // Bottom of kernel stack for this process
71     enum procstate state;   // Process state
72     volatile int pid;       // Process ID
73     struct proc *parent;    // Parent process
74     struct trapframe *tf;   // Trap frame for current syscall
75     struct context *context; // swtch() here to run process
76     void *chan;             // If non-zero, sleeping on chan
77     int killed;             // If non-zero, have been killed
78     struct file *ofile[NOFILE]; // Open files
79     struct inode *cwd;      // Current directory
80     char name[16];          // Process name (debugging)
81
82     // PS shit
83     _Bool inuse; // If it's being run by a CPU or not
84     int ticks; // How many ticks has accumulated
85     // Lottery shit
86     int tickets;
87
88 };
89

```

proc.h [+]

74,1-4

79%

I store some of the information from pstat in proc, to make the code cleaner. Inuse is effectively (state==RUNNING)

```

96 // Let's allow the process table to be public
97 struct ptable_type {
98     struct spinlock lock;
99     struct proc proc[NPROC];
100 };
101 extern struct ptable_type ptable;
102
103
104 void setproctickets(struct proc* pp, int n);

```

I mentioned previously that I moved the definition of ptable's struct to proc.h so I could access ptable from other files

Sysfunc.h

```
25 25 int sys_uptime(void);  
26 +int sys_getpinfo(void);  
27 +int sys_settickets(void);  
26 28
```

syscall.c

```
105 105 [SYS_uptime] sys_uptime,  
106 +[SYS_getpinfo] sys_getpinfo,  
107 +[SYS_settickets] sys_settickets,  
106 108 };
```

Sysproc.c

```
11 int sys_settickets(void)  
12 {  
13     int number_of_tickets;  
14     // Error  
15     if(argint(0, &number_of_tickets) < 0)  
16         return -1;  
17  
18     acquire(&ptable.lock);  
19     setproctickets(proc, number_of_tickets);  
20     release(&ptable.lock);  
21  
22     return 0;  
23 }
```

Sys_settickets changes the number of tickets of a process, and returns 0 on success and -1 on error.

```

25 int sys_getpinfo(void)
26 {
27     acquire(&ptable.lock);
28     struct pstat* target;
29     if(argint(0, (int*)&target)) < 0)
30         return -1;
31
32     for(struct proc* p=ptable.proc; p != &(ptable.proc[NPROC]); p++)
33     {
34         const int index = p - ptable.proc;
35         if(p->state != UNUSED)
36         {
37             target->pid[index] = p->pid;
38             target->ticks[index] = p->ticks;
39             target->tickets[index] = p->tickets;
40             target->inuse[index] = p->inuse;
41             target->state[index] = p->state;
42         }
43     }
44     target->total_tickets = total_tickets;
45     release(&ptable.lock);
46     return 0;
47 }
48 }

```

Sys_getpinfo fills a pstat struct with data from the ptable. I have to do this all within a lock for the case of multiple CPUs

xv6/user

Ps.c

```
1 /* This utility runs unit tests*/
2 #include "types.h"
3 #include "pstat.h"
4 #include "user.h"
5
6 int ps(){
7     // Get process info first
8     struct pstat pinfo = {0};
9     if(-1 == getpinfo(&pinfo)){
10         return 0;
11         fprintf(1, "\n\t FAILURE\n");
12     }
13     // Our process id
14     const int current_pid = getpid();
15     // Total ticks of not sleeping processes(not including this one)
16     int total_ticks = 0;
17     // Where we are in pinfo
18     int current_entry = 0;
19     for(int i=0;i<NPROC;i++){
20         {
21             if(pinfo.pid[i] == current_pid)
22             {
23                 current_entry = i;
24             }else if(pinfo.state[i] != SLEEPING)
```

ps.c 24,3-9 Top

```

23         current_entry = i;
24     }else if(pinfo.state[i] != SLEEPING)
25     {
26         total_ticks += pinfo.ticks[i];
27     }
28 }
29
30
31 // Header
32 fprintf(stdout,"This process: %d\n",
33         getpid());
34 fprintf(stdout,"Total tickets: %d\n\n",
35         pinfo.total_tickets);
36 fprintf(stdout,
37         "PID\tTicks\tTickets\tState\tE%\tA%\n");
38
39 // Body
40 for(int i=0;i<NPROC;i++){
41     if(pinfo.pid[i]==0)continue;
42     _Bool skip_yield = 0;
43     // What we expect the percentage of ticks to be
44     float expected_yield =
45         100*

```



```

44     float expected_yield =
45         100*
46         (float)pinfo.tickets[i]/
47         ((float)pinfo.total_tickets - pinfo
48 .tickets[current_entry]);
49     // What the actual percentage of ticks
is
49     float actual_yield =
50         100*
51         (float)pinfo.ticks[i]/
52         (float)total_ticks;
53
54     if(pinfo.state[i] == 2 || pinfo.pid[i]
55 == current_pid)
56     {
57         skip_yield = 1;
58     }
59     int ey_left =
60         (int)expected_yield;
61     int ey_right =
62         (int)((expected_yield-ey_left)*10);
63     int ay_left =
64         (int)actual_yield;
65     int ay_right =
66         (int)((actual_yield-ay_left)*10);

```

ps.c

44,3-9

57%


~~~~~

86,5-17

Bot

This is the user program “PS” that displays the PID, ticks, tickets, state, expected ticks %, and actual ticks %, of each process, as well as whether the process is in use by displaying “|”(not in use) or “>”(in use)

## Bm.c

```
1  /*
2  * BM - Baby Maker
3  * Makes an arbitrary amount of babies, letting
4  * them live for
5  * eternity */
6  #include "types.h"
7  #include "user.h"
8  int main(int argc, char** argv)
9  {
10     if(argc<2){
11         fprintf(stdout, "Usage: bm child_1_tickets [child_2_tickets]...\n");
12         exit();
13     }
14     fprintf(stdout, "Mother %d created\n", getpid());
15     for(int i=1;i<argc;i++){
16         const int pid = fork();
17         if(pid<0){
18             fprintf(stderr, "Stillbirth Occurred\n");
19             exit();
20         }
21         if(!pid)
22             {
23                 const int t = atoi(argv[i]); // number of tickets
24                 settickets(t);
25                 fprintf(stdout, "Child %d created with %d tickets\n", getpid(), t);
26                 // Loop forever
27                 while(1);
28                 fprintf(stdout, "Child %d exiting\n", getpid());
29                 exit();
30             }
31     }
32     for(int i=1;i<argc;i++){
33         wait();
34     }
35     fprintf(stdout, "Parent exiting\n");
36     exit();
37 }
```

bm.c 1,1 Top bm.c 37,1 Bot

This program creates an arbitrary number of child processes of a specified number of tickets

## Grapher.c

```

1 #include "types.h"
2 #include "fcntl.h"
3 #include "user.h"
4 #include "pstat.h"
5 /*
6  * This programs monitors two processes and records their numer of ticks every
7  * so and so intervals
8  */
9 // Sampling frequency
10 const int SAMP_PERIOD = 75;
11 // Total time sampled
12 const int SAMP_WINDOW = 2000;
13
14 // How many tickets the grapher has
15 const int GRAPHER_PRIORITY = 1000;
16 // Where we're writing to
17 const char*const OUTPUT_FILE = "graph.csv";
18
19 int main(int argc, char** argv)
20 {
21     settickets(GRAPHER_PRIORITY);
22
23     // How many processes we're observing
24     const int PROCESS_QUANTITY = argc - 1;
25
26     // list of pids
27     int processes[PROCESS_QUANTITY];
28     for(int i=0;i<PROCESS_QUANTITY;i++)
29         processes[i] = atoi(argv[i+1]);
30
31     // Let's get the pstat struct
32     struct pstat pinfo = {0};
33     if(0>getpinfo(&pinfo))
34     {
35         fprintf(stderr, "getpinfo() failed\n");
36         exit();
37     }
38
39     // Open file

```

```

40 //unlink(OUTPUT_FILE);
41 int fp = stdout; //open(OUTPUT_FILE, O_WRONLY);
42
43 // Write header
44 fprintf(fp, "time, ");
45 for(int i=0; i<PROCESS_QUANTITY; i++)
46 {
47     fprintf(fp, "%d, ", processes[i]);
48 }
49 write(fp, "\n", 1);
50
51 // List of indexes in the pstat struct
52 // that give us the processes we want
53 int pindices[PROCESS_QUANTITY];
54 for(int index=0; index<NPROC; index++)
55 {
56     for(int i=0; i<PROCESS_QUANTITY; i++)
57     {
58         if(pinfo.pid[index] == processes[i])
59         {
60             pindices[i] = index;
61         }
62     }
63 }
64
65
66
67
68 int time_passed = 0;
69 while(1)
70 {
71     // Update pinfo
72     getpinfo(&pinfo);
73
74     fprintf(fp, "%d, ", uptime());
75
76     // Fill ticks
77     for(int i=0; i<PROCESS_QUANTITY; i++)
78     {
79

```

```

57     {
58         if(pinfo.pid[index] == processes[i])
59         {
60             pindices[i] = index;
61         }
62     }
63 }
64
65
66
67
68 int time_passed = 0;
69 while(1)
70 {
71     // Update pinfo
72     getpinfo(&pinfo);
73
74     fprintf(fp, "%d, ", uptime());
75
76     // Fill ticks
77     for(int i=0; i<PROCESS_QUANTITY;i++)
78     {
79
80         fprintf(fp, "%d, ", pinfo.ticks[pindices[i]]);
81     }
82     write(fp, "\n", 1);
83
84     // End if needed
85     if(time_passed>=SAMP_WINDOW)
86         break;
87     // Sleep
88     sleep(SAMP_PERIOD);
89     time_passed+=SAMP_PERIOD;
90 }
91
92 close(fp);
93
94
95 exit();
96 }

```

grapher.c

57,3-9

Bot

This is a program that creates a graph plot in the form of a CSV file which is then fed into Octave to create the graph proper.

## Makefile.mk

|    |    |                 |
|----|----|-----------------|
| 2  | 2  | # user programs |
| 3  | 3  | USER_PROGS := \ |
| 4  | 4  | cat\            |
|    | 5  | + bf\           |
|    | 6  | + bm\           |
| 5  | 7  | echo\           |
| 6  | 8  | forktest\       |
| 7  | 9  | grep\           |
|    | 10 | + grapher\      |
| 8  | 11 | init\           |
| 9  | 12 | kill\           |
| 10 | 13 | ln\             |
| 11 | 14 | ls\             |
| 12 | 15 | mkdir\          |
|    | 16 | + ps\           |
| 13 | 17 | rm\             |
| 14 | 18 | sh\             |
| 15 | 19 | stressfs\       |

Included here are the three programs I mentioned, plus a Brainf\*ck interpreter (for fun)

## Ulib.c

|     |     |     |                                                            |          |
|-----|-----|-----|------------------------------------------------------------|----------|
|     |     | ✱   | @@ -4,6 +4,13 @@                                           |          |
| 4   | 4   |     | #include                                                   | "user.h" |
| 5   | 5   |     | #include                                                   | "x86.h"  |
| 6   | 6   |     |                                                            |          |
|     |     | 7   | +char                                                      |          |
|     |     | 8   | +getchar()                                                 |          |
|     |     | 9   | +{                                                         |          |
|     |     | 10  | +     char c;                                              |          |
|     |     | 11  | +     int cc = read(stdin, &c, 1);                         |          |
|     |     | 12  | +     return cc== -1?-1:c;                                 |          |
|     |     | 13  | +}                                                         |          |
| 7   | 14  |     | char*                                                      |          |
| 8   | 15  |     | strcpy(char *s, char *t)                                   |          |
| 9   | 16  |     | {                                                          |          |
|     |     | ✱   | @@ -24,7 +31,7 @@ strcmp(const char *p, const char *q)     |          |
| 24  | 31  |     | }                                                          |          |
| 25  | 32  |     |                                                            |          |
| 26  | 33  |     | uint                                                       |          |
| 27  |     |     | -strlen(char *s)                                           |          |
|     | 34  |     | +strlen(const char *s)                                     |          |
| 28  | 35  |     | {                                                          |          |
| 29  | 36  |     | int n;                                                     |          |
| 30  | 37  |     |                                                            |          |
|     |     | ✱   | @@ -103,3 +110,8 @@ memmove(void *vdst, void *vsrc, int n) |          |
| 103 | 110 |     | *dst++ = *src++;                                           |          |
| 104 | 111 |     | return vdst;                                               |          |
| 105 | 112 |     | }                                                          |          |
|     |     | 113 | +                                                          |          |
|     |     | 114 | +// Constants                                              |          |
|     |     | 115 | +const int stdin = 0;                                      |          |
|     |     | 116 | +const int stdout = 1;                                     |          |
|     |     | 117 | +const int stderr = 2;                                     |          |

Added the function "getchar" and made a few function accept constant pointers



## User.h

|    |    |                                                        |
|----|----|--------------------------------------------------------|
| 3  | 3  | <code>+#include "pstat.h"</code>                       |
| 3  | 4  |                                                        |
| 4  | 5  | <code>struct stat;</code>                              |
| 5  | 6  |                                                        |
| ✱  |    | <code>@@ -8,14 +9,14 @@ int fork(void);</code>         |
| 8  | 9  | <code>int exit(void) __attribute__((noreturn));</code> |
| 9  | 10 | <code>int wait(void);</code>                           |
| 10 | 11 | <code>int pipe(int*);</code>                           |
| 11 |    | <code>-int write(int, void*, int);</code>              |
| 12 |    | <code>+int write(int, const void*, int);</code>        |
| 12 | 13 | <code>int read(int, void*, int);</code>                |
| 13 | 14 | <code>int close(int);</code>                           |
| 14 | 15 | <code>int kill(int);</code>                            |
| 15 | 16 | <code>int exec(char*, char**);</code>                  |
| 16 |    | <code>-int open(char*, int);</code>                    |
| 17 |    | <code>+int open(const char*, int);</code>              |
| 17 | 18 | <code>int mknod(char*, short, short);</code>           |
| 18 |    | <code>-int unlink(char*);</code>                       |
| 19 |    | <code>+int unlink(const char*);</code>                 |
| 19 | 20 | <code>int fstat(int fd, struct stat*);</code>          |
| 20 | 21 | <code>int link(char*, char*);</code>                   |
| 21 | 22 | <code>int mkdir(char*);</code>                         |
| ✱  |    | <code>@@ -25,20 +26,30 @@ int getpid(void);</code>     |
| 25 | 26 | <code>char* sbrk(int);</code>                          |
| 26 | 27 | <code>int sleep(int);</code>                           |
| 27 | 28 | <code>int uptime(void);</code>                         |
| 29 |    | <code>+int getpinfo(struct pstat*);</code>             |
| 30 |    | <code>+int settickets(int);</code>                     |

```

29 32 // user library functions (ulib.c)
30 33 int stat(char*, struct stat*);
31 34 char* strcpy(char*, char*);
32 35 void *memmove(void*, void*, int);
33 36 char* strchr(const char*, char c);
34 37 int strcmp(const char*, const char*);
35 -void printf(int, char*, ...);
38 +void fprintf(int, const char*, ...);
39 +void putchar(char c);
40 +char getchar(void);
36 41 char* gets(char*, int max);
37 -uint strlen(char*);
42 +uint strlen(const char*);
38 43 void* memset(void*, int, uint);
39 44 void* malloc(uint);
45 +void* calloc(uint, uint);
40 46 void free(void*);
41 47 int atoi(const char*);
42 48
49 +// Constants
50 +extern const int stdin;
51 +extern const int stdout;
52 +extern const int stderr;
53 +

```

Here, I added my two syscalls as well as some personalization  
**usys.S**

```

29 29 SYSCALL(sbrk)
30 30 SYSCALL(sleep)
31 31 SYSCALL(uptime)
32 +SYSCALL(getpinfo)
33 +SYSCALL(settickets)

```

Added syscalls to assembly file


## Results

```
$ ps
This process: 8
Total tickets: 400

PID      Ticks    Tickets  State   E%     A%
|1        1         1        2       -      -
|2        1        100      2       -      -
|5       952       100      3      33.3   32.8
|4        1        100      2       -      -
|6     1944       200      3      66.6   67.1
>8        0        100      4       -      -
$
```

The result of running `ps` sometime after `bm 100 200`

Note that while the default for any process spawning from init(PID 1) is 1 ticket, but the default for any process spawning from the shell(PID 2) is 100, because the shell's ticket is manually set in the user program `sh`.



This graph shows the accumulated ticks of two processes PID5 and PID6 with 100 and 200 tickets, respectively.